

SAI VENKATA NAGENDRA ABHAY RAJ GURRAM

Jersey City, NJ, United States | +1 815-618-1777 | abhayrajgurram13@gmail.com
linkedin.com/in/abhaygurram | github.com/Abhayraj1307

PROFESSIONAL SUMMARY

Machine Learning and Data Science graduate student specializing in AI-driven decision systems, NLP, and data-intensive applications. Experienced in building end-to-end ML pipelines, LLM-based systems, and optimization frameworks that generate actionable insights from structured and unstructured data. Strong focus on scalable system design, real-time decision workflows, and solving ambiguous problems with measurable outcomes.

TECHNICAL SKILLS

Programming & Querying: Python, SQL

Data & Statistics: Pandas, NumPy, Exploratory Data Analysis (EDA), Statistical Analysis, Hypothesis Testing, A/B Testing, Experimental Design

Machine Learning: Regression, Classification, Clustering, Feature Engineering, Model Evaluation (Precision, Recall, F1), Cross-Validation

Deep Learning & NLP: PyTorch, Transformers, BERT, Hugging Face, Tokenization, Embeddings

LLM & AI Systems: Retrieval-Augmented Generation (RAG), Prompt Engineering, Multi-Agent Systems, Tool Integration, Decision Intelligence Systems

Data Engineering & BI: ETL Pipelines, Data Processing, Data Modeling, SQL Joins, Window Functions

Cloud & Deployment: AWS (S3, SageMaker), FastAPI, Streamlit, Model Deployment, API Development, Scalable System Design

Tools & Systems: Git, Version Control, Debugging, Modular Architecture

PROJECT EXPERIENCE

NovaMind – AI Decision Intelligence Platform

- Built an AI system integrating LLM reasoning, RAG pipelines, and external tools to generate structured decisions from unstructured inputs
- Designed multi-agent architecture to decompose workflows into analysis, reasoning, and execution layers
- Implemented FAISS-based vector retrieval to improve contextual relevance
- Developed Streamlit interface for real-time decision outputs
- Result: Reduced manual analysis effort by ~25–30% in test scenarios

AI Multi-Cloud GPU Cost Optimization & Orchestration System

- Built ETL pipelines to ingest and normalize GPU pricing data across cloud providers
- Developed decision engine to evaluate cost-performance trade-offs for ML workloads
- Modeled workload scenarios to simulate infrastructure decisions
- Designed API-based system for real-time deployment recommendations
- Result: Achieved ~20% lower estimated cost compared to baseline scenarios

Rationale Agent – Incident Intelligence System

- Built LLM-based system to analyze incident data and generate root cause hypotheses
- Implemented reasoning workflows for hypothesis comparison and evidence extraction
- Designed UI for structured, decision-ready outputs
- Result: Reduced incident triage time by ~35–40% in test cases

Exo-Cortex Debugger – AI-Powered Code Debugging Tool

- Built Python execution sandbox to run code and capture runtime errors
- Implemented traceback parsing and LLM-based debugging explanations
- Generated structured insights and fix recommendations
- Result: Reduced debugging time for common errors by ~25–30%

BERT-Based Sentiment Analysis System

- Fine-tuned transformer models for sentiment classification
- Performed preprocessing, tokenization, and embedding generation
- Built dashboards to visualize sentiment insights
- Result: Achieved strong classification performance with consistent accuracy

PROFESSIONAL EXPERIENCE

Quality Assurance Intern

Laurus Labs, Visakhapatnam, India | Jan 2022 – Jun 2022

- Performed batch record review and data validation to ensure regulatory compliance
- Identified deviations and supported corrective and preventive actions (CAPA)
- Worked with structured datasets in compliance-driven workflows
- Result: Improved documentation accuracy and reduced validation errors

EDUCATION

Stevens Institute of Technology

Master's in Business Analytics & Artificial Intelligence | May 2026

GPA: 3.51

GITAM Deemed University

Bachelor's in Pharmacy

CERTIFICATIONS

Apple Search Ads Certification

Bank of America Investment Banking Simulation