

SAI VENKATA NAGENDRA ABHAY RAJ GURRAM

Jersey City, NJ | 815-618-1777 | nagendraabhaygurram@gmail.com

www.linkedin.com/in/abhaygurram | github.com/Abhayraj1307

SUMMARY

Graduate student in Business Intelligence & Analytics (AI focus) building production-style AI tools: an LLM incident intelligence agent, a Python debugging copilot, BERT+RAG sentiment systems, and F1 race strategy models. Strong in Python, SQL, ML, and LLMs, with end-to-end experience from data prep to deployment. Actively seeking a Summer 2025 internship in Data Science, Machine Learning, or AI Engineering.

EDUCATION

Stevens Institute of Technology, Hoboken, NJ

M.S. in Business Analytics & AI.

Expected May 2026

GPA: 3.25 / 4.00 | Coursework: NLP & Deep Learning, Web Mining, Applied Statistics, Data Visualization, Business Analytics

Gandhi Institute of Technology and Management (GITAM), Visakhapatnam, India

Bachelor of Pharmacy

2019 - 2023

TECHNICAL PROJECTS

Rationale Agent – Incident Intelligence | Python, Streamlit, LLM APIs, Markdown Parsing

GitHub: github.com/Abhayraj1307/rationale-agent-incident-intel

- Built an LLM-powered incident assistant that ingests markdown post-mortems and infers the dominant root-cause category plus competing alternatives.
- Designed a decision summary layer that aggregates evidence, computes a heuristic confidence score, and groups claims by cause with impacted systems and supporting snippets for faster RCA reviews.

Exo-Cortex – Python Debugging Copilot | Python, Streamlit, Sandbox Execution, Unit Tests

GitHub: github.com/Abhayraj1307/exocortex-debugger

- Developed an AI-driven debugger where users paste Python code, run it in a sandbox, capture traceback, and receive LLM-generated explanations and corrected code suggestions.
- Modularized the backend (sandbox runner, traceback parser, LLM client) and added targeted unit tests, making the tool safer to extend without breaking the debugging flow.

Twitter Sentiment Analysis with BERT & RAG | Python, Hugging Face, Streamlit

- Fine-tuned a BERT model on tweet data and combined it with a Retrieval-Augmented Generation layer to output both sentiment labels and natural-language rationales.
- Deployed the model as an interactive Streamlit app and on Hugging Face, achieving ~12% accuracy improvement over a baseline sentiment classifier.

F1 Race Dynamics – Predictive Modeling for Podium Finishes | XGBoost, SHAP, Tableau

- Engineered features from lap-by-lap race data (tire strategy, pit windows, safety car events) and trained gradient-boosted models to predict podium finishes.
- Used SHAP to identify tire strategy and pit timing as key predictors (~82% validation accuracy) and built a Tableau dashboard to explore race strategy scenarios.

INDUSTRY EXPERIENCE

Laurus Labs Ltd., Visakhapatnam, India

Quality Assurance Intern, May 2022 – June 2022

- Worked with batch records, stability data, and deviation reports for large-scale pharma production, improving clarity and consistency of QA documentation.
- Helped standardize basic QA checks and data workflows for review, reducing rework and making internal audits easier for the core team.

SKILLS

Programming: Python, SQL, R

ML & NLP: scikit-learn, XGBoost, LightGBM, SHAP, Transformers (BERT), spaCy, clustering, text classification, predictive modeling

LLMs & Agents: Prompt engineering, RAG pipelines, LLM-based tools for debugging and incident intelligence, Hugging Face Transformers

Data Apps & Visualization: Streamlit, Tableau, Power BI, Plotly

Data & Infra: PostgreSQL, SQL Server, Docker, FastAPI, Git/GitHub, dbt (basic), Airflow (basic), AWS SageMaker, GCP Cloud Run

CERTIFICATIONS & PUBLICATION

Apple Ads Search Certification | Nanotechnology: A Maker's Course | Drug Discovery | Drug Development

Publication: UV Analysis of Ritonavir